

Paper B-09: Knowledge Graphs: A Constraint-Driven Flux Dynamics (CDFD) Perspective

Steve Bico Mujjabi, MD

Independent Researcher

Founder, Vura Labs

Kampala, Uganda

ORCID: 0009-0001-0556-5516

August 2026

Shared Part IV notation and evidence boundaries are defined in `PART_IV_METHODS_AND_CLAIM_BOUNDARY.md`. This paper states only its local observables, comparison, and failure conditions.

Abstract

This paper develops a falsifiable CDFD/AFL model of Knowledge Graphs. It maps queries, facts, relations, and update streams to driving flux Φ , schema, compute, provenance, consistency, and quality limits to active constraint C , inference, curation, indexing, schema evolution, and conflict resolution to responsiveness S , and accumulated ontology, provenance, stale facts, and design decisions to structural memory M_s . The target outcome is retrieval, consistency, coverage, and update reliability, tested against query logs, graph topology, provenance, edit history, and benchmark tasks. The paper treats universality as a shared model architecture rather than an assertion that unlike systems have the same material mechanism. It states the negative result that would make the mapping unnecessary and places the domain claim inside the cumulative CDFD programme and its reproducible runtime boundary.

1 Introduction

A knowledge graph is not merely a database; it is a topological map of human understanding. Entities (nodes) are linked by relationships (edges), allowing a machine to infer context.

However, knowledge is not static. New discoveries continuously rewrite the relationships between entities. If a knowledge graph possesses a rigid, unyielding schema, it will rapidly become obsolete. In CDFD terms, a static database is a thermodynamically dead system lacking the elasticity to process new informational flux.

2 The Semantic Adaptive Surface

We govern the informational topology using the Adaptive Operating Ratio:

$$\Psi_s = \left(\frac{\Phi}{C} \right) \cdot S \cdot M_s \quad (1)$$

For a large-scale knowledge graph (e.g., a global medical database):

- **Flux** (Φ): The continuous stream of new data ingestion, user queries, and logical inference sweeps across the network.
- **Constraint** (C): The strictness of the ontology schema, data-type validation rules, and the computational cost of multi-hop graph traversal.
- **Responsiveness** (S): The ability of the graph to dynamically generate new edge types, infer latent relationships, and actively reshape its own ontology to fit incoming data.
- **Memory** (M_s): Hardcoded edge weights and legacy taxonomy rules. High memory locks the graph into obsolete worldviews, violently resisting the integration of paradigm-shifting information.

2.1 Topological Sclerosis and Dynamic Inference

In a perfectly elastic knowledge graph, high S and low M_s allow the system to continuously adapt. When a novel piece of data arrives (a spike in Φ), the graph effortlessly creates a new edge, absorbing the flux without hitting a constraint ($\Psi_s \approx 1$).

If the graph’s ontology is highly rigid (high M_s , high C), the incoming novel data (Φ) violates the strict schema validation. The data is rejected. The flow of truth into the system is bottlenecked ($\Psi_s \ll 1$). The knowledge graph enters a state of topological sclerosis. It ceases to accurately reflect reality because its internal mathematical memory refuses to adapt to external constraints.

To survive, a true semantic engine must utilize machine learning to constantly prune obsolete edges and fluidly relax its constraints C , ensuring that the informational topology perfectly mirrors the shifting thermodynamic reality it attempts to model.

3 Measurement Layer

4 Cross-Domain Model Upgrade: Mechanism, Scale, and Tests

4.1 Observable Translation

The paper-specific translation begins with an observable chain rather than a verbal analogy. For Knowledge Graphs, the proposed drive is queries, facts, relations, and update streams. The active limitation is schema, compute, provenance, consistency, and quality limits. Responsiveness is carried by inference, curation, indexing, schema evolution, and conflict resolution,

while retained state is represented by accumulated ontology, provenance, stale facts, and design decisions. The outcome to explain is retrieval, consistency, coverage, and update reliability.

Table 1: Operational CDFD/AFL map for Paper B-09.

Term	Paper-specific observable meaning	Measurement question
Φ	queries, facts, relations, and update streams	What rate, load, gradient, or demand enters the focal system?
C	schema, compute, provenance, consistency, and quality limits	Which measured bottleneck limits transfer, function, or recovery?
S	inference, curation, indexing, schema evolution, and conflict resolution	Which process changes effective capacity or reroutes the load?
M_s	accumulated ontology, provenance, stale facts, and design decisions	Which retained state makes equal present loads produce different futures?
Y	retrieval, consistency, coverage, and update reliability	Which outcome is predicted out of sample, with uncertainty?




Figure 1: Paper B-09 observable CDFD/AFL architecture for Knowledge Graphs. Solid arrows give the proposed forward chain; dashed arrows make history-dependent feedback explicit.

4.2 Paper-Specific Causal Hypothesis

For this paper, the hypothesized causal sequence is specific. A change in queries, facts, relations, and update streams arrives faster than inference, curation, indexing, schema evolution, and conflict resolution can compensate. This raises or redistributes schema, compute, provenance, consistency, and quality limits. If the load persists, the retained state encoded by accumulated ontology, provenance, stale facts, and design decisions changes the next response, so a later event of equal magnitude can produce a different trajectory in retrieval, consistency, coverage, and update reliability. The scientific task is to estimate that sequence against the best ordinary model of the domain, not merely to redescribe the outcome after it occurs.

4.3 Cross-Scale Translation Without Mechanistic Erasure

For the engineered systems family, the cross-domain translation compares demand, designed capacity, feedback control, dependency topology, and accumulated technical state. The strongest engineering use is prospective: the mapping must identify a bottleneck, a control action, or a recovery signature before the failure occurs. [2, 4, 3, 1] The required scale declaration for this

paper is milliseconds to decades; triples to knowledge ecosystems. At short scales, the key question is whether the incoming drive is transmitted, buffered, or diverted. At intermediate scales, response changes effective capacity and network exposure. At long scales, retained structure can alter the state space itself. A result is genuinely cross-scale only when the aggregation rule is stated and the same conclusion is not an artifact of averaging.

4.4 Evidence Ladder and Discriminating Tests

Table 2: Evidence ladder for Paper B-09. Each rung can fail independently.

Test	Design	Discriminating result
Proxy audit	Measure query logs, graph topology, provenance, edit history, and benchmark tasks and preregister how each record maps to Φ , C , S , M_s , and Y .	The mapping fails if proxies are non-identifiable, circular, or change meaning across cases.
Load ramp	Apply or observe a rapid update or contradiction introduced into a dense sub-graph while estimating the dominant constraint and response time.	A threshold claim requires a reproducible nonlinear change and uncertainty bounds, not a visually chosen breakpoint.
Recovery and memory	Match present load across units with different histories, then compare recovery paths.	The memory claim requires history-dependent divergence after current conditions are controlled.
Model comparison	Compare the baseline domain model with and without the CDFD/AFL state variables using held-out data.	The paper’s stated falsifier is: constraint-memory terms do not improve error propagation or repair prediction.

4.5 Research Programme and Boundary Conditions

The immediate research programme has four steps. First, construct a documented dataset from query logs, graph topology, provenance, edit history, and benchmark tasks and retain raw units before normalization. Second, fit a baseline domain model and report its calibration and out-of-sample error. Third, add the smallest identifiable CDFD/AFL state extension and test whether it improves prediction of retrieval, consistency, coverage, and update reliability. Fourth, repeat the analysis across the declared scale range and publish negative cases as part of the cross-domain translation audit.

5 Conclusion

Data is a fluid; the database schema is the pipe. By applying CDFD to knowledge graph architecture, we frame the hypothesis that semantic reasoning engines must operate as living, biological-like tissues. They must actively mutate their own structural memory (M_s) and relax their constraints (C) to maintain the optimal $\Psi_s \approx 1$ flow of truth.

Conflict of Interest

The author declares no conflict of interest.

Data Availability

Part-specific runtime diagnostics are under each Part's `outputs/` directory.

Release-wide code: `Part_E_Synthesis/supplementary/`.

Generated summaries: `Part_E_Synthesis/outputs/`.

Figures: `Part_E_Synthesis/figures/`.

6 Reference Layer

References

- [1] Per Bak, Chao Tang, and Kurt Wiesenfeld. Self-organized criticality: An explanation of the $1/f$ noise. *Physical Review Letters*, 59(4):381–384, 1987. DOI: <https://doi.org/10.1103/PhysRevLett.59.381>.
- [2] Albert-Laszlo Barabasi and Reka Albert. Emergence of scaling in random networks. *Science*, 286(5439):509–512, 1999. DOI: <https://doi.org/10.1126/science.286.5439.509>.
- [3] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948. DOI: <https://doi.org/10.1002/j.1538-7305.1948.tb01338.x>.
- [4] Duncan J. Watts and Steven H. Strogatz. Collective dynamics of small-world networks. *Nature*, 393:440–442, 1998. DOI: <https://doi.org/10.1038/30918>.