

Paper B-09: Knowledge Graphs: A Constraint-Driven Flux Dynamics (CDFD) Perspective

Steve Bico Mujjabi, MD

Independent Researcher

Founder, Vura Labs

Kampala, Uganda

ORCID: 0009-0001-0556-5516

May 2026

Abstract

This paper develops a falsifiable CDFD/AFL model of Knowledge Graphs. It maps queries, facts, relations, and update streams to driving flux Φ , schema, compute, provenance, consistency, and quality limits to active constraint C , inference, curation, indexing, schema evolution, and conflict resolution to responsiveness S , and accumulated ontology, provenance, stale facts, and design decisions to structural memory M_s . The target outcome is retrieval, consistency, coverage, and update reliability, tested against query logs, graph topology, provenance, edit history, and benchmark tasks. The paper treats universality as a shared model architecture rather than an assertion that unlike systems have the same material mechanism. It states the negative result that would make the mapping unnecessary and places the domain claim inside the cumulative CDFD programme and its reproducible runtime boundary.

Symbol Conventions

CDFD/AFL Part IV symbol convention.

Symbol	Part IV meaning	Cross-domain reading
Φ	Driving flux or throughput	Energy, matter, signal, capital, information, traffic, force, or biological load entering a system
C	Active constraint or capacity burden	Bottleneck, resistance, boundary, scarcity, toxicity, regulation, topology, latency, or load-bearing limit
S	Responsiveness or adaptive surface	The system's ability to reconfigure, buffer, route, repair, learn, or preserve function under changing load
M_s	Structural memory	Stored damage, hysteresis, inherited topology, institutional memory, trained weights, ecological legacy, or retained state
Ψ_s	Adaptive operating ratio $(\Phi/C) \cdot S \cdot M_s$	Local bookkeeping gauge for constrained operation, overload, recovery, or regime change; not a universal constant
Λ	Life Number or capstone surplus ratio where explicitly defined	Reserved for Part II-derived life-number notation or synthesis-level surplus measures, never an undefined decoration

Greek-symbol convention.

Symbol	Local role	Interpretation rule
α	Accumulation or reinforcement coefficient	Sets how fast a pathway, constraint, habit, cascade, or retained structure grows under load
β	Relaxation, decay, clearance, or washout coefficient	Sets loss, recovery, damping, forgetting, degradation, or return toward baseline
γ	Coupling, diffusion, redistribution, or transfer coefficient	Sets spread across nodes, layers, tissues, markets, infrastructures, media, or physical phases
δ, ϵ	Perturbation, tolerance, small offset, or numerical guard	Used only as local thresholds, stability margins, measurement tolerances, or stress increments
η, λ, μ, ν	Efficiency, rate, baseline, intensity, or frequency terms	Meaning is fixed by the equation, subscript, and domain-specific proxy in the paragraph where the term appears
ρ, σ, τ, ϕ	Density/correlation, variability/noise, time-scale, phase, or field terms	Local coefficients only; subscripts bind them to the named pathway, domain, or experiment
Ω, Λ	Domain/state space; Life Number where defined	Ω names a modeled state space; Λ carries life-number or synthesis notation only when the paper defines it

1 Series Context

Part I sets the flow-constraint language used here [3]. Part II develops the Life Number and the origins-of-life use of the same notation [4]. Part III applies AFL to biology and medicine

[5]. Part IV [6], DOI <https://doi.org/10.5281/zenodo.20395901>, is the cross-domain stress test: it asks where the notation clarifies measurement, and where it becomes only a metaphor.

2 Introduction

2.1 The Muijabi Universal Capacity Law

Building directly upon the Fundamental Physics (Part I) [3], the **Muijabi Universal Capacity Law** proposes that across all scales—from black hole event horizons to global financial markets and power grids—driving high flux (Φ) against a rigid topological constraint (C) can drive system saturation. When Φ/C approaches critical thresholds, the system does not scale linearly; it undergoes a topological phase transition.

2.2 The Muijabi Universal Memory Law

The **Muijabi Universal Memory Law** frames persistent systemic collapse. When a complex network fails to clear its structural memory (M_s) and its relaxation parameter decays ($\tau_{relax} \rightarrow \infty$), temporary stressors become permanent rigidities. This is the proposed CDFD mechanism behind ecological death, economic depressions, and infrastructural gridlock.

A knowledge graph is not merely a database; it is a topological map of human understanding. Entities (nodes) are linked by relationships (edges), allowing a machine to infer context.

However, knowledge is not static. New discoveries continuously rewrite the relationships between entities. If a knowledge graph possesses a rigid, unyielding schema, it will rapidly become obsolete. In CDFD terms, a static database is a thermodynamically dead system lacking the elasticity to process new informational flux.

3 The Semantic Adaptive Surface

We govern the informational topology using the Universal Adaptive Ratio:

$$\Psi_s = \left(\frac{\Phi}{C} \right) \cdot S \cdot M_s \quad (1)$$

For a large-scale knowledge graph (e.g., a global medical database):

- **Flux** (Φ): The continuous stream of new data ingestion, user queries, and logical inference sweeps across the network.
- **Constraint** (C): The strictness of the ontology schema, data-type validation rules, and the computational cost of multi-hop graph traversal.
- **Responsiveness** (S): The ability of the graph to dynamically generate new edge types, infer latent relationships, and actively reshape its own ontology to fit incoming data.
- **Memory** (M_s): Hardcoded edge weights and legacy taxonomy rules. High memory locks the graph into obsolete worldviews, violently resisting the integration of paradigm-shifting information.

3.1 Topological Sclerosis and Dynamic Inference

In a perfectly elastic knowledge graph, high S and low M_s allow the system to continuously adapt. When a novel piece of data arrives (a spike in Φ), the graph effortlessly creates a new edge, absorbing the flux without hitting a constraint ($\Psi_s \approx 1$).

If the graph’s ontology is highly rigid (high M_s , high C), the incoming novel data (Φ) violates the strict schema validation. The data is rejected. The flow of truth into the system is bottlenecked ($\Psi_s \ll 1$). The knowledge graph enters a state of topological sclerosis. It ceases to accurately reflect reality because its internal mathematical memory refuses to adapt to external constraints.

To survive, a true semantic engine must utilize machine learning to constantly prune obsolete edges and fluidly relax its constraints C , ensuring that the informational topology perfectly mirrors the shifting thermodynamic reality it attempts to model.

4 Measurement Layer

The first job in any domain is to choose proxies before drawing conclusions. Φ may be a rate of material, energy, information, capital, or signal flow. C is the bottleneck that slows or redirects that flow. S measures the system’s ability to reconfigure under stress, and M_s records the past load that still changes present behavior. A useful application gives those four terms observable definitions, a time scale, and a failure condition. If a standard domain model explains the same data more cleanly, the CDFD/AFL mapping should give way.

4.1 Current Runtime Diagnostic: Universal Network Cascade

The release-local discovery script keeps this paper tied to the current CDFD Runtime [7], including RK4 field updates, surface dynamics, structural-memory diffusion, and a finite-value audit. In the 2026-06-04 runtime pass, the localized hub-drive stress test ended with hub $\Psi_s = 5.21$, peak network $\Psi_s = 5.21$, 21 nodes above $\Psi_s > 1.5$, and 37 maximum memory-locked nodes. I treat those numbers as a runtime sanity check and as a target for later domain measurement, not as a substitute for field data.

5 Major Universal Upgrade: Mechanism, Scale, and Tests

5.1 Observable Translation

The paper-specific translation begins with an observable chain rather than a verbal analogy. For Knowledge Graphs, the proposed drive is queries, facts, relations, and update streams. The active limitation is schema, compute, provenance, consistency, and quality limits. Responsiveness is carried by inference, curation, indexing, schema evolution, and conflict resolution, while retained state is represented by accumulated ontology, provenance, stale facts, and design decisions. The outcome to explain is retrieval, consistency, coverage, and update reliability.

Table 1: Operational CDFD/AFL map for Paper B-09.

Term	Paper-specific observable meaning	Measurement question
Φ	queries, facts, relations, and update streams	What rate, load, gradient, or demand enters the focal system?
C	schema, compute, provenance, consistency, and quality limits	Which measured bottleneck limits transfer, function, or recovery?
S	inference, curation, indexing, schema evolution, and conflict resolution	Which process changes effective capacity or reroutes the load?
M_s	accumulated ontology, provenance, stale facts, and design decisions	Which retained state makes equal present loads produce different futures?
Y	retrieval, consistency, coverage, and update reliability	Which outcome is predicted out of sample, with uncertainty?




Figure 1: Paper B-09 observable CDFD/AFL architecture for Knowledge Graphs. Solid arrows give the proposed forward chain; dashed arrows make history-dependent feedback explicit.

5.2 Minimal Coupled Model

A paper-level model must separate throughput, constraint accumulation, response, and history. One deliberately minimal form is

$$Y(t) = \frac{\Phi(t)S(t)}{\epsilon + C(t)}, \quad (2)$$

$$\frac{dC}{dt} = \alpha\Phi(t) - \beta S(t)C(t) + \gamma\mathcal{L}C(t), \quad (3)$$

$$\frac{dM_s}{dt} = \eta C(t) - \mu M_s(t), \quad (4)$$

$$\Psi_s(t) = \frac{\hat{\Phi}(t)}{\epsilon + \hat{C}(t)} \hat{S}(t) [1 + \omega \hat{M}_s(t)]. \quad (5)$$

Here Y is the measured output, $\epsilon > 0$ prevents a singular denominator, $\mathcal{L}$ is a spatial or network coupling operator, and hats denote explicitly normalized variables. The sign and magnitude of ω must be estimated: memory can preserve useful organization or amplify damage. No fixed threshold is universal before the variables, normalization, and observation scale are declared.

For this paper, the hypothesized causal sequence is specific. A change in queries, facts, relations, and update streams arrives faster than inference, curation, indexing, schema evolution, and conflict resolution can compensate. This raises or redistributes schema, compute, provenance, consistency, and quality limits. If the load persists, the retained state encoded by

accumulated ontology, provenance, stale facts, and design decisions changes the next response, so a later event of equal magnitude can produce a different trajectory in retrieval, consistency, coverage, and update reliability. The scientific task is to estimate that sequence against the best ordinary model of the domain, not merely to redescribe the outcome after it occurs.

5.3 Cross-Scale Universality Without Mechanistic Erasure

For the engineered systems family, the universal comparison is between demand, designed capacity, feedback control, dependency topology, and accumulated technical state. The strongest engineering use is prospective: the mapping must identify a bottleneck, a control action, or a recovery signature before the failure occurs. [2, 9, 8, 1] The CDFD programme is cumulative: Part I supplies the flow-constraint foundation [3]; Part II asks when maintained throughput supports persistent organized states [4]; Part III develops adaptive limitation and biological memory [5]; and Part IV [6] tests how much of that architecture survives translation. The CDFD Runtime [7] supplies a reproducible numerical stress surface, but it does not substitute for the domain evidence named here.

The required scale declaration for this paper is milliseconds to decades; triples to knowledge ecosystems. At short scales, the key question is whether the incoming drive is transmitted, buffered, or diverted. At intermediate scales, response changes effective capacity and network exposure. At long scales, retained structure can alter the state space itself. A result is genuinely cross-scale only when the aggregation rule is stated and the same conclusion is not an artifact of averaging.

5.4 Evidence Ladder and Discriminating Tests

Table 2: Evidence ladder for Paper B-09. Each rung can fail independently.

Test	Design	Discriminating result
Proxy audit	Measure query logs, graph topology, provenance, edit history, and benchmark tasks and preregister how each record maps to Φ , C , S , M_s , and Y .	The mapping fails if proxies are non-identifiable, circular, or change meaning across cases.
Load ramp	Apply or observe a rapid update or contradiction introduced into a dense subgraph while estimating the dominant constraint and response time.	A threshold claim requires a reproducible nonlinear change and uncertainty bounds, not a visually chosen breakpoint.
Recovery and memory	Match present load across units with different histories, then compare recovery paths.	The memory claim requires history-dependent divergence after current conditions are controlled.
Model comparison	Compare the baseline domain model with and without the CDFD/AFL state variables using held-out data.	The paper’s stated falsifier is: constraint-memory terms do not improve error propagation or repair prediction.

The strongest practical design combines a prospective perturbation with a recovery window. The load-ramp phase estimates where constraint begins to rise faster than output. The recovery phase estimates β and μ , separating fast relaxation from persistent memory. A topology or coupling intervention then asks whether rerouting changes the cascade as predicted. Null

results are informative because they identify domains in which the universal notation adds no measurable value.

5.5 Research Programme and Boundary Conditions

The immediate research programme has four steps. First, construct a documented dataset from query logs, graph topology, provenance, edit history, and benchmark tasks and retain raw units before normalization. Second, fit a baseline domain model and report its calibration and out-of-sample error. Third, add the smallest identifiable CDFD/AFL state extension and test whether it improves prediction of retrieval, consistency, coverage, and update reliability. Fourth, repeat the analysis across the declared scale range and publish negative cases as part of the universality audit.

Three boundaries are non-negotiable. Correlation between flow and failure does not identify a constraint mechanism. A fitted memory term does not prove that the stored state has the same material meaning in another domain. Finally, a runtime cascade is evidence about the implementation, not evidence that the real system follows it. The paper becomes stronger when it can say precisely where the translation stops.

6 Conclusion

Data is a fluid; the database schema is the pipe. By applying CDFD to knowledge graph architecture, we frame the hypothesis that semantic reasoning engines must operate as living, biological-like tissues. They must actively mutate their own structural memory (M_s) and relax their constraints (C) to maintain the optimal $\Psi_s \approx 1$ flow of truth.

Conflict of Interest

The author declares no conflict of interest.

Data Availability

Part-specific runtime diagnostics are under each Part's `outputs/` directory.

Release-wide code: `Part_E_Synthesis/supplementary/`.

Generated summaries: `Part_E_Synthesis/outputs/`.

Figures: `Part_E_Synthesis/figures/`.

7 Reference Layer

References

- [1] Per Bak, Chao Tang, and Kurt Wiesenfeld. Self-organized criticality: An explanation of the $1/f$ noise. *Physical Review Letters*, 59(4):381–384, 1987. DOI: <https://doi.org/10.1103/PhysRevLett.59.381>.

- [2] Albert-Laszlo Barabasi and Reka Albert. Emergence of scaling in random networks. *Science*, 286(5439):509–512, 1999. DOI: <https://doi.org/10.1126/science.286.5439.509>.
- [3] Steve Bico Mujjabi. Cdfd part i: Fundamental physics - constraint-driven field theory and flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20250821>.
- [4] Steve Bico Mujjabi. Cdfd part ii: Origins of life and tri-regime bioenergetics - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20264779>.
- [5] Steve Bico Mujjabi. Cdfd part iii: Adaptive flux limitation biology and medicine - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20344695>.
- [6] Steve Bico Mujjabi. Cdfd part iv: Universal adaptive flux limitation - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20395901>.
- [7] Steve Bico Mujjabi. Cdfd runtime: Constraint-driven flux dynamics and cdfd execution engine, 2026. DOI: <https://doi.org/10.5281/zenodo.20343160>.
- [8] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948. DOI: <https://doi.org/10.1002/j.1538-7305.1948.tb01338.x>.
- [9] Duncan J. Watts and Steven H. Strogatz. Collective dynamics of small-world networks. *Nature*, 393:440–442, 1998. DOI: <https://doi.org/10.1038/30918>.