

Paper B-09: Knowledge Graphs: A Constraint-Driven Flux Dynamics (CDFD) Perspective

Steve Bico Mujjabi, MD

Independent Researcher

ORCID: 0009-0001-0556-5516

May 2026

Abstract

This paper asks whether Constraint-Driven Flux Dynamics (CDFD) and Adaptive Flux Limitation (AFL) can give the named system a sharper audit language. The working variables are driving flux Φ , constraint C , surface responsiveness S , and structural memory M_s . The useful question is plain: what can be measured, what would count as overload, and what result would make the mapping fail?

1 Series Context

Part I sets the flow–constraint language used here [1]. Part II develops the Life Number and the origins-of-life use of the same notation [2]. Part III applies AFL to biology and medicine [3]. Part IV is the cross-domain stress test: it asks where the notation clarifies measurement, and where it becomes only a metaphor.

2 Introduction

Notation. In Part IV, C names the active constraint or capacity burden in the system being discussed. Φ is the driving flow, S is the surface response, and M_s is retained structural memory. Λ appears only where the Part II Life Number is being referenced [2]. Other Greek symbols are local coefficients whose meaning is set by the equation in which they appear.

2.1 The Mujjabi Universal Capacity Law

Building directly upon the Fundamental Physics (Part I) [1], the **Mujjabi Universal Capacity Law** proposes that across all scales—from black hole event horizons to global financial markets and power grids—driving high flux (Φ) against a rigid topological constraint (C) can drive system saturation. When Φ/C approaches critical thresholds, the system does not scale linearly; it undergoes a topological phase transition.

2.2 The Mujjabi Universal Memory Law

The **Mujjabi Universal Memory Law** frames persistent systemic collapse. When a complex network fails to clear its structural memory (M_s) and its relaxation parameter decays ($\tau_{relax} \rightarrow \infty$), temporary stressors become permanent rigidities. This is the proposed CDFD mechanism behind ecological death, economic depressions, and infrastructural gridlock.

A knowledge graph is not merely a database; it is a topological map of human understanding. Entities (nodes) are linked by relationships (edges), allowing a machine to infer context.

However, knowledge is not static. New discoveries continuously rewrite the relationships between entities. If a knowledge graph possesses a rigid, unyielding schema, it will rapidly become obsolete. In CDFD terms, a static database is a thermodynamically dead system lacking the elasticity to process new informational flux.

3 The Semantic Adaptive Surface

We govern the informational topology using the Universal Adaptive Ratio:

$$\Psi_s = \left(\frac{\Phi}{C} \right) \cdot S \cdot M_s \quad (1)$$

For a large-scale knowledge graph (e.g., a global medical database):

- **Flux** (Φ): The continuous stream of new data ingestion, user queries, and logical inference sweeps across the network.
- **Constraint** (C): The strictness of the ontology schema, data-type validation rules, and the computational cost of multi-hop graph traversal.
- **Responsiveness** (S): The ability of the graph to dynamically generate new edge types, infer latent relationships, and actively reshape its own ontology to fit incoming data.
- **Memory** (M_s): Hardcoded edge weights and legacy taxonomy rules. High memory locks the graph into obsolete worldviews, violently resisting the integration of paradigm-shifting information.

3.1 Topological Sclerosis and Dynamic Inference

In a perfectly elastic knowledge graph, high S and low M_s allow the system to continuously adapt. When a novel piece of data arrives (a spike in Φ), the graph effortlessly creates a new edge, absorbing the flux without hitting a constraint ($\Psi_s \approx 1$).

If the graph's ontology is highly rigid (high M_s , high C), the incoming novel data (Φ) violates the strict schema validation. The data is rejected. The flow of truth into the system is bottlenecked ($\Psi_s \ll 1$). The knowledge graph enters a state of topological sclerosis. It ceases to accurately reflect reality because its internal mathematical memory refuses to adapt to external constraints.

To survive, a true semantic engine must utilize machine learning to constantly prune obsolete edges and fluidly relax its constraints C , ensuring that the informational topology perfectly mirrors the shifting thermodynamic reality it attempts to model.

4 Measurement Layer

The first job in any domain is to choose proxies before drawing conclusions. Φ may be a rate of material, energy, information, capital, or signal flow. C is the bottleneck that slows or redirects that flow. S measures the system’s ability to reconfigure under stress, and M_s records the past load that still changes present behavior. A useful application gives those four terms observable definitions, a time scale, and a failure condition. If a standard domain model explains the same data more cleanly, the CDFD/AFL mapping should give way.

4.1 Current Runtime Diagnostic: Universal Network Cascade

The diagnostic script `Part_E_Synthesis/supplementary/run_partiv_discovery.py` keeps this paper tied to the current CDFD Runtime [4], including RK4 field updates, surface dynamics, structural-memory diffusion, and a finite-value audit. In the May 2026 run, the localized hub-drive stress test ended with hub $\Psi_s = 14.8$, peak network $\Psi_s = 14.8$, 21 nodes above $\Psi_s > 1.5$, and 37 maximum memory-locked nodes. I treat those numbers as a runtime sanity check and as a target for later domain measurement, not as a substitute for field data.

5 Conclusion

Data is a fluid; the database schema is the pipe. By applying CDFD to knowledge graph architecture, we frame the hypothesis that semantic reasoning engines must operate as living, biological-like tissues. They must actively mutate their own structural memory (M_s) and relax their constraints (C) to maintain the optimal $\Psi_s \approx 1$ flow of truth.

Conflict of Interest

The author declares no conflict of interest.

Data Availability

Release-local runtime diagnostics are available under each Part’s `outputs/` folder, with release-wide synthesis artifacts under `Part_E_Synthesis/supplementary/`, `Part_E_Synthesis/outputs/`, and `Part_E_Synthesis/figures/`.

6 Reference Layer

References

- [1] Steve Bico Mujjabi. Cdfd part i: Fundamental physics - constraint-driven field theory and flux dynamics, 2026.
- [2] Steve Bico Mujjabi. Cdfd part ii: Origins of life and tri-regime bioenergetics - constraint-driven flux dynamics, 2026.
- [3] Steve Bico Mujjabi. Cdfd part iii: Adaptive flux limitation biology and medicine - constraint-driven flux dynamics, 2026.
- [4] Steve Bico Mujjabi. Cdfd runtime: Constraint-driven flux dynamics and cdfl execution engine, 2026.