

Paper D-09: AI Neural Computation: A Constraint-Driven Flux Dynamics (CDFD) Perspective

Steve Bico Mujjabi, MD

Independent Researcher

Founder, Vura Labs

Kampala, Uganda

ORCID: 0009-0001-0556-5516

May 2026

Abstract

This paper treats AI neural computation as a CDFD/AFL domain-application hypothesis. Token flow, activation flow, gradient updates, memory retrieval, tool calls, and inference demand are read as candidate fluxes Φ moving through constraints C such as context length, compute budget, memory bandwidth, latency, alignment filters, and training-data structure. The goal is to define measurable overload and recovery variables, not to claim that an AI system has life, intention, or consciousness.

Symbol Conventions

CDFD/AFL Part IV symbol convention.

Symbol	Part IV meaning	Cross-domain reading
Φ	Driving flux or throughput	Energy, matter, signal, capital, information, traffic, force, or biological load entering a system
C	Active constraint or capacity burden	Bottleneck, resistance, boundary, scarcity, toxicity, regulation, topology, latency, or load-bearing limit
S	Responsiveness or adaptive surface	The system's ability to reconfigure, buffer, route, repair, learn, or preserve function under changing load
M_s	Structural memory	Stored damage, hysteresis, inherited topology, institutional memory, trained weights, ecological legacy, or retained state
Ψ_s	Adaptive operating ratio $(\Phi/C) \cdot S \cdot M_s$	Local bookkeeping gauge for constrained operation, overload, recovery, or regime change; not a universal constant
Λ	Life Number or capstone surplus ratio where explicitly defined	Reserved for Part II-derived life-number notation or synthesis-level surplus measures, never an undefined decoration

Greek-symbol convention.

Symbol	Local role	Interpretation rule
α	Accumulation or reinforcement coefficient	Sets how fast a pathway, constraint, habit, cascade, or retained structure grows under load
β	Relaxation, decay, clearance, or washout coefficient	Sets loss, recovery, damping, forgetting, degradation, or return toward baseline
γ	Coupling, diffusion, redistribution, or transfer coefficient	Sets spread across nodes, layers, tissues, markets, infrastructures, media, or physical phases
δ, ϵ	Perturbation, tolerance, small offset, or numerical guard	Used only as local thresholds, stability margins, measurement tolerances, or stress increments
η, λ, μ, ν	Efficiency, rate, baseline, intensity, or frequency terms	Meaning is fixed by the equation, subscript, and domain-specific proxy in the paragraph where the term appears
ρ, σ, τ, ϕ	Density/correlation, variability/noise, time-scale, phase, or field terms	Local coefficients only; subscripts bind them to the named pathway, domain, or experiment
Ω, Λ	Domain/state space; Life Number where defined	Ω names a modeled state space; Λ carries life-number or synthesis notation only when the paper defines it

1 Series Context

Part I sets the flow-constraint language used here [3]. Part II develops the Life Number and the origins-of-life use of the same notation [4]. Part III applies AFL to biology and medicine [5]. This paper uses those references for notation and series continuity,

while Part IV [6], DOI <https://doi.org/10.5281/zenodo.20395901>, supplies the cross-domain stress-test archive; it does not transfer biological claims onto AI systems.

2 AI Computation as Constrained Flow

Modern AI systems move information through layered constraints. A prompt pushes tokens into a context window; attention routes signals through a trained parameter field; retrieval systems add external memory; inference hardware sets latency, bandwidth, and energy limits; safety filters and tool policies add additional routing constraints. CDFD/AFL reads the system as a flow problem with memory: outputs depend on current demand and on structure laid down during training, fine-tuning, caching, and retrieval.

3 Candidate Variable Map

CDFD term	AI-system reading	Example proxy
Φ	inference or update flow	token rate, request load, gradient norm
C	active constraint	context length, memory bandwidth, latency, policy filter
S	responsiveness	routing change, retrieval update, tool selection, scaling
M_s	structural memory	weights, embeddings, cache, retrieval corpus, history
Ψ_s	operating ratio	stable, saturated, brittle, or over-loaded regime

4 Failure Modes

AI overload does not need mystical language. It can mean context saturation, retrieval conflict, latency collapse, hallucination under sparse grounding, memory contamination, brittle tool routing, or training-inference mismatch. Those states are measurable only when the system boundary is clear. A local model, a hosted model, an agent loop, and a retrieval-augmented system have different Φ , C , S , and M_s proxies.

5 Measurement Layer

Useful measurements include token throughput, context occupancy, retrieval hit rate, cache age, latency, refusal/error rate, grounding score, tool-call success, memory bandwidth, gradient stability, and energy cost. A CDFD/AFL analysis must name the failure condition before interpreting Ψ_s . Without that condition, the notation only renames ordinary engineering metrics.

6 Current Runtime Diagnostic

The Part B adapter sweep includes an artificial-intelligence row, and the Part E cascade diagnostic checks finite behavior and memory locking in the current CDFD Runtime [7]. These outputs are release evidence for the runtime, not evidence that any deployed AI system has been validated by this paper.

7 Tests That Can Break the Mapping

The AI mapping weakens if the proposed variables cannot predict saturation, latency, grounding failure, or recovery better than existing system metrics; if M_s cannot be separated from ordinary model state; or if the analysis depends on anthropomorphic language instead of measurable system behavior.

8 Major Universal Upgrade: Mechanism, Scale, and Tests

8.1 Observable Translation

The paper-specific translation begins with an observable chain rather than a verbal analogy. For AI Neural Computation, the proposed drive is training data, gradients, context, and inference requests. The active limitation is compute, memory, data quality, architecture, bandwidth, and evaluation. Responsiveness is carried by optimization, routing, retrieval, tool use, and oversight, while retained state is represented by weights, optimizer state, context, fine-tuning, and deployment feedback. The outcome to explain is loss, capability, calibration, robustness, and reliability.

Table 1: Operational CDFD/AFL map for Paper D-09.

Term	Paper-specific observable meaning	Measurement question
Φ	training data, gradients, context, and inference requests	What rate, load, gradient, or demand enters the focal system?
C	compute, memory, data quality, architecture, bandwidth, and evaluation	Which measured bottleneck limits transfer, function, or recovery?
S	optimization, routing, retrieval, tool use, and oversight	Which process changes effective capacity or reroutes the load?
M_s	weights, optimizer state, context, fine-tuning, and deployment feedback	Which retained state makes equal present loads produce different futures?
Y	loss, capability, calibration, robustness, and reliability	Which outcome is predicted out of sample, with uncertainty?




Figure 1: Paper D-09 observable CDFD/AFL architecture for AI Neural Computation. Solid arrows give the proposed forward chain; dashed arrows make history-dependent feedback explicit.

8.2 Minimal Coupled Model

A paper-level model must separate throughput, constraint accumulation, response, and history. One deliberately minimal form is

$$Y(t) = \frac{\Phi(t)S(t)}{\epsilon + C(t)}, \quad (1)$$

$$\frac{dC}{dt} = \alpha\Phi(t) - \beta S(t)C(t) + \gamma\mathcal{L}C(t), \quad (2)$$

$$\frac{dM_s}{dt} = \eta C(t) - \mu M_s(t), \quad (3)$$

$$\Psi_s(t) = \frac{\widehat{\Phi}(t)}{\epsilon + \widehat{C}(t)} \widehat{S}(t) [1 + \omega \widehat{M}_s(t)]. \quad (4)$$

Here Y is the measured output, $\epsilon > 0$ prevents a singular denominator, $\mathcal{L}$ is a spatial or network coupling operator, and hats denote explicitly normalized variables. The sign and magnitude of ω must be estimated: memory can preserve useful organization or amplify damage. No fixed threshold is universal before the variables, normalization, and observation scale are declared.

For this paper, the hypothesized causal sequence is specific. A change in training data, gradients, context, and inference requests arrives faster than optimization, routing, retrieval, tool use, and oversight can compensate. This raises or redistributes compute, memory, data quality, architecture, bandwidth, and evaluation. If the load persists, the retained state encoded by weights, optimizer state, context, fine-tuning, and deployment feedback changes the next response, so a later event of equal magnitude can produce a different trajectory in loss, capability, calibration, robustness, and reliability. The scientific task is to estimate that sequence against the best ordinary model of the domain, not merely to redescribe the outcome after it occurs.

8.3 Cross-Scale Universality Without Mechanistic Erasure

For the domain applications family, the universal comparison is a disciplined translation test across systems with very different material mechanisms. A valid mapping preserves

the baseline science of the domain, declares the scale of every variable, and earns its place through a discriminating prediction rather than resemblance alone. [2, 9, 8, 1] The CDFD programme is cumulative: Part I supplies the flow–constraint foundation [3]; Part II asks when maintained throughput supports persistent organized states [4]; Part III develops adaptive limitation and biological memory [5]; and Part IV [6] tests how much of that architecture survives translation. The CDFD Runtime [7] supplies a reproducible numerical stress surface, but it does not substitute for the domain evidence named here.

The required scale declaration for this paper is nanoseconds to deployment lifetimes; parameters to AI ecosystems. At short scales, the key question is whether the incoming drive is transmitted, buffered, or diverted. At intermediate scales, response changes effective capacity and network exposure. At long scales, retained structure can alter the state space itself. A result is genuinely cross-scale only when the aggregation rule is stated and the same conclusion is not an artifact of averaging.

8.4 Evidence Ladder and Discriminating Tests

Table 2: Evidence ladder for Paper D-09. Each rung can fail independently.

Test	Design	Discriminating result
Proxy audit	Measure training curves, ablations, activations, evaluations, incidents, and distribution shifts and preregister how each record maps to Φ , C , S , M_s , and Y .	The mapping fails if proxies are non-identifiable, circular, or change meaning across cases.
Load ramp	Apply or observe a capacity bottleneck or distribution shift after training while estimating the dominant constraint and response time.	A threshold claim requires a reproducible nonlinear change and uncertainty bounds, not a visually chosen breakpoint.
Recovery and memory	Match present load across units with different histories, then compare recovery paths.	The memory claim requires history-dependent divergence after current conditions are controlled.
Model comparison	Compare the baseline domain model with and without the CDFD/AFL state variables using held-out data.	The paper’s stated falsifier is: CDFD variables do not improve scaling, adaptation, or failure prediction.

The strongest practical design combines a prospective perturbation with a recovery window. The load-ramp phase estimates where constraint begins to rise faster than output. The recovery phase estimates β and μ , separating fast relaxation from persistent memory. A topology or coupling intervention then asks whether rerouting changes the cascade as predicted. Null results are informative because they identify domains in which the universal notation adds no measurable value.

8.5 Research Programme and Boundary Conditions

The immediate research programme has four steps. First, construct a documented dataset from training curves, ablations, activations, evaluations, incidents, and distribution shifts and retain raw units before normalization. Second, fit a baseline domain model and report its calibration and out-of-sample error. Third, add the smallest identifiable CDFD/AFL

state extension and test whether it improves prediction of loss, capability, calibration, robustness, and reliability. Fourth, repeat the analysis across the declared scale range and publish negative cases as part of the universality audit.

Three boundaries are non-negotiable. Correlation between flow and failure does not identify a constraint mechanism. A fitted memory term does not prove that the stored state has the same material meaning in another domain. Finally, a runtime cascade is evidence about the implementation, not evidence that the real system follows it. The paper becomes stronger when it can say precisely where the translation stops.

9 Conclusion

AI neural computation is a strong stress test for Part IV because it has visible flow, constraint, responsiveness, and memory. The claim stays modest: CDFD/AFL may help audit overload and recovery in AI systems when the variables are measured plainly.

Conflict of Interest

The author declares no conflict of interest.

Data Availability

Part-specific runtime diagnostics are under each Part's `outputs/` directory.

Release-wide code: `Part_E_Synthesis/supplementary/`.

Generated summaries: `Part_E_Synthesis/outputs/`.

Figures: `Part_E_Synthesis/figures/`.

10 Reference Layer

References

- [1] Per Bak, Chao Tang, and Kurt Wiesenfeld. Self-organized criticality: An explanation of the $1/f$ noise. *Physical Review Letters*, 59(4):381–384, 1987. DOI: <https://doi.org/10.1103/PhysRevLett.59.381>.
- [2] C. S. Holling. Resilience and stability of ecological systems. *Annual Review of Ecology and Systematics*, 4:1–23, 1973. DOI: <https://doi.org/10.1146/annurev.es.04.110173.000245>.
- [3] Steve Bico Mujjabi. Cdfd part i: Fundamental physics - constraint-driven field theory and flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20250821>.
- [4] Steve Bico Mujjabi. Cdfd part ii: Origins of life and tri-regime bioenergetics - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20264779>.

- [5] Steve Bico Mujjabi. Cdfd part iii: Adaptive flux limitation biology and medicine - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20344695>.
- [6] Steve Bico Mujjabi. Cdfd part iv: Universal adaptive flux limitation - constraint-driven flux dynamics, 2026. DOI: <https://doi.org/10.5281/zenodo.20395901>.
- [7] Steve Bico Mujjabi. Cdfd runtime: Constraint-driven flux dynamics and cdf execution engine, 2026. DOI: <https://doi.org/10.5281/zenodo.20343160>.
- [8] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27(3):379–423, 1948. DOI: <https://doi.org/10.1002/j.1538-7305.1948.tb01338.x>.
- [9] A. M. Turing. The chemical basis of morphogenesis. *Philosophical Transactions of the Royal Society of London. Series B*, 237(641):37–72, 1952. DOI: <https://doi.org/10.1098/rstb.1952.0012>.